

The Political Implications of Machine-Verified Mathematics: What the Cathedral Means for Policy

Jason Robert Gochanour

April 19, 2026

Abstract

The Cathedral is a machine-verified formal proof system that reduces the Riemann Hypothesis to seven axioms in Lean 4, built through human-AI collaboration in 23 days. This paper examines the political and policy implications of five aspects of the project: (1) the verified instability certificate capability and its relevance to critical infrastructure protection, (2) the human-AI collaboration model and its implications for AI governance, (3) machine verification as a tool for evidence-based policymaking, (4) the geopolitics of mathematical capability and open science, and (5) the question of responsible disclosure timing. We argue that the Cathedral’s principal policy impact is not its mathematical content (which does not break cryptography or create weapons) but its *methodology*: machine-verified reasoning applied to complex systems. This methodology, if adopted by policymakers, could transform how governments evaluate scientific claims, certify critical infrastructure, and regulate AI systems.

Contents

1	Scope and Caveats	2
2	Critical Infrastructure and Verified Vulnerability Analysis	2
2.1	The Capability	2
2.2	Policy Implications	3
3	AI Governance: A Concrete Case Study	3
3.1	What the Cathedral Demonstrates	3
3.2	Policy Implications for AI Regulation	4
4	Machine Verification and Evidence-Based Policy	4
4.1	The Problem of Contested Science	4
4.2	What Machine Verification Offers	5
4.3	Feasibility	5
5	Geopolitics of Mathematical Capability	6
5.1	The Landscape	6
5.2	Open Science as Defense	6
5.3	The One-Person Lab	6

6	Responsible Disclosure and Timing	7
6.1	The Decision Framework	7
6.2	Recommended Sequence	7
7	The Verification Imperative	8
8	Conclusion	8

1 Scope and Caveats

This paper is written for a policy audience. It assumes no mathematical background. It does not advocate for any political position. Its purpose is to map the political landscape that the Cathedral’s existence creates, so that informed decisions can be made about dissemination and application.

Three important caveats:

1. **The Cathedral does not prove the Riemann Hypothesis.** It reduces RH to two well-understood axioms. The policy implications come from the *methodology*, not the mathematical result.
2. **The Cathedral does not break cryptography.** A companion paper (*cathedral-security.tex*) establishes this rigorously. The mathematical tools used (sieves, Parseval’s theorem, spectral analysis) provide no computational shortcuts for factoring or discrete logarithms.
3. **The Cathedral does not create weapons.** The companion paper (*cathedral-dualuse.tex*) identifies dual-use risks, but none involve weapons of mass destruction or capabilities beyond what existing mathematical tools already provide.

With these caveats, the political implications are real but subtle. They concern *how we make decisions about complex systems*, not what we can blow up.

2 Critical Infrastructure and Verified Vulnerability Analysis

2.1 The Capability

The companion paper *cathedral-dualuse.tex* identifies a novel capability: **verified instability certificates**. Using the Cathedral’s proof methodology, it is possible in principle to produce a machine-checked proof that a specific sequence of events will cause a critical infrastructure system (e.g., a power grid) to fail.

This is qualitatively different from simulation-based vulnerability analysis:

	Simulation	Verified Certificate
Correctness guarantee	Probabilistic	Mathematical
Can miss edge cases	Yes	No (within model)
Requires expertise to evaluate	High	Low (compiler checks)
Reproducible	Sometimes	Always
Falsifiable	Yes	No (within axioms)

2.2 Policy Implications

1. **Defensive priority.** If verified instability certificates are possible, grid operators and government agencies should develop them defensively—finding vulnerabilities before adversaries do. The mathematical tools are publicly available; the defense must stay ahead of potential offense.
2. **Classification question.** Should verified instability certificates for specific grid configurations be classified? Current CEII (Critical Energy Infrastructure Information) regulations protect grid topology data but do not address mathematical proofs about grid behavior. Policy may need to be updated.
3. **International dimension.** Any nation with graduate students in mathematics and access to Lean 4 could develop this capability. It is not a technology that can be controlled through export restrictions. The defense must be institutional (deploy verified analysis defensively) rather than technological (restrict access to tools).
4. **NERC/FERC implications.** The North American Electric Reliability Corporation (NERC) and Federal Energy Regulatory Commission (FERC) may need to update Critical Infrastructure Protection (CIP) standards to address formally verified vulnerability analysis.

3 AI Governance: A Concrete Case Study

3.1 What the Cathedral Demonstrates

The Cathedral provides one of the most detailed case studies of human-AI collaboration on a complex intellectual task. The companion paper *cathedral-ai.tex* documents the division of labor:

- The AI generated 85% of code by volume.
- The AI proposed proof strategies that were adopted.
- The AI designed axiom elimination campaigns.
- The human provided architectural vision, quality control, and mathematical taste.
- The Lean compiler served as the ultimate verification layer.

3.2 Policy Implications for AI Regulation

1. **The “Amplifier” Model.** The Cathedral supports the view that current AI systems are *capability amplifiers* rather than autonomous agents. The AI could not have designed the Cathedral alone; the human could not have built it in 23 days alone. Regulatory frameworks should account for this collaborative dynamic rather than treating AI as either a tool or an agent.
2. **Verification as Safety.** The Cathedral demonstrates that formal verification (Lean 4) can serve as a trust anchor for AI-generated output. The human does not need to check every line of AI-generated code—the compiler does that. This “verification-as-safety” model could be applied to other domains where AI generates complex outputs:
 - AI-generated legal documents verified by formal logic.
 - AI-generated engineering designs verified by formal methods.
 - AI-generated scientific claims verified by machine-checked proofs.
3. **Attribution and Accountability.** The Cathedral raises the question: who is responsible for AI-assisted work? The project credits the AI as a collaborative contributor, not merely a tool. This has implications for:
 - Patent law (can AI be a co-inventor?).
 - Academic authorship (can AI be a co-author?).
 - Liability (if AI-generated code causes harm, who is responsible?).
4. **Democratization of Expertise.** A single person with AI assistance produced a 78-file formal verification project in 23 days—work that would traditionally require a research team of 5–10 over months. This has implications for:
 - Science funding models (smaller grants, faster timelines).
 - Workforce policy (the solo researcher + AI as a viable research unit).
 - International competitiveness (nations that enable AI-assisted research gain a multiplier).

4 Machine Verification and Evidence-Based Policy

4.1 The Problem of Contested Science

Policy decisions increasingly depend on complex scientific claims that are difficult for non-experts to evaluate:

- Climate models with thousands of parameters.
- Epidemiological models (COVID-19 response).
- Economic forecasts.

- Drug efficacy claims.
- Infrastructure safety assessments.

In each case, policy relies on *trusting experts*, which creates vulnerability to:

- Genuine scientific disagreement (different models, different conclusions).
- Bad faith manipulation (models tuned to support a predetermined conclusion).
- Communication failure (accurate science, misunderstood by policymakers).

4.2 What Machine Verification Offers

The Cathedral demonstrates a model where:

1. Every assumption is **named** (axiom registry).
2. Every deduction is **machine-checked** (Lean compiler).
3. The dependency graph is **auditable** (`#print axioms`).
4. The entire argument is **reproducible** (deterministic compilation).

Political translation: A machine-verified policy analysis would look like this:

“We conclude that Policy X will reduce carbon emissions by $15\% \pm 3\%$, conditional on:

- Axiom 1: GDP growth remains below 3%.
- Axiom 2: No new coal plants are built.
- Axiom 3: Carbon capture technology achieves 90% efficiency.

Every other step in the analysis is machine-verified. If you disagree with our conclusion, you must disagree with one of these three axioms. The compiler has checked everything else.”

This would transform policy debates from “I don’t trust your model” to “I disagree with Axiom 2.” The argument shifts from authority to assumptions—a healthier basis for democratic deliberation.

4.3 Feasibility

Near-term: Machine verification of simple policy models (tax revenue projections, population forecasts) is achievable with current tools.

Medium-term: Machine verification of engineering safety analyses (bridge loads, dam stability, building codes) would require formalized physics libraries in Lean 4 but is structurally feasible.

Long-term: Machine verification of climate models, epidemiological models, and economic forecasts would require significant investment in formalized domain-specific libraries, but the Cathedral shows that 20,000 lines of verified mathematics is achievable in weeks, not years.

5 Geopolitics of Mathematical Capability

5.1 The Landscape

Formal verification capability is currently concentrated in a small number of institutions:

- **Lean/Mathlib:** Primarily academic, open-source, centered at Carnegie Mellon, Imperial College, and the broader Lean community.
- **AlphaProof:** Google DeepMind, proprietary.
- **LeanDojo/Lean Copilot:** Academic, open-source.
- **National labs:** DARPA HACMS program (verified software for military systems); various classified programs.

The Cathedral demonstrates that a single individual with commercial AI tools can produce significant formal verification work. This lowers the barrier to entry for *any* nation or non-state actor.

5.2 Open Science as Defense

The Cathedral’s papers argue that open publication is the responsible choice because:

1. The mathematical tools are already public (textbook material).
2. Defenders benefit more from understanding the tools than attackers (asymmetric advantage to the defender).
3. Suppressing publication does not deny capability; it only denies *awareness of the capability* to defenders.

However, the *timing* of publication matters. Responsible disclosure in cybersecurity follows a protocol: notify the vendor, allow time to patch, then publish. The same principle could apply to mathematical infrastructure analysis:

1. Develop verified vulnerability analysis internally.
2. Share with infrastructure operators (NERC, grid operators).
3. Allow time for remediation.
4. Publish methodology (not specific vulnerabilities).

5.3 The One-Person Lab

The Cathedral’s most politically significant implication may be the simplest: a single person with AI assistance can now do work that previously required institutional resources. This means:

- Traditional gatekeepers (universities, national labs, research councils) no longer control access to frontier mathematical research.
- The cost of significant formal verification has dropped from “research team $\times$ years” to “one person $\times$ weeks.”
- Small nations, independent researchers, and non-state actors can produce verified mathematical analysis at a level previously reserved for elite institutions.

This is neither good nor bad. It is a structural change in who has access to mathematical power. Policy should acknowledge this reality rather than pretend it can be reversed.

6 Responsible Disclosure and Timing

6.1 The Decision Framework

The decision of when and how to make the Cathedral public involves:

1. **Mathematical content:** The proof architecture itself. This is not sensitive—it concerns the logical structure of a well-known conjecture.
2. **Engineering applications:** The companion papers on engineering, energy, and futures propose novel applications. These are speculative but grounded, and early disclosure to relevant industries allows defensive development.
3. **Dual-use analysis:** The dark mirror paper identifies capabilities that, while not new in principle, are made more precise by the Cathedral’s methodology. Early sharing with infrastructure security professionals allows defensive preparation.
4. **Methodology:** The human-AI collaboration model, the axiom management system, and the verification architecture are the most broadly useful contributions. These should be published openly, as they benefit the entire verification community.

6.2 Recommended Sequence

1. **First:** Share with trusted colleagues in relevant fields for technical review and risk assessment.
2. **Second:** Share the dual-use analysis with infrastructure security professionals (CISA, NERC, relevant national agencies).
3. **Third:** Publish the mathematical and methodological papers openly (math, CS, philosophy, AI, Lean, foundations, engineering).
4. **Fourth:** Publish the dual-use and energy papers after infrastructure operators have had time to evaluate the analysis.
5. **Throughout:** The proof code itself can be published at any stage, as it does not contain operational vulnerabilities.

7 The Verification Imperative

The deepest political implication of the Cathedral is not about infrastructure, AI, or cryptography. It is about **truth**.

In a political environment characterized by contested facts, algorithmic misinformation, and declining trust in institutions, the Cathedral offers a model of *mechanically verified truth*. The Lean compiler does not care about politics, funding, nationality, or ideology. It checks whether each logical step follows from the previous one. That is all it does. And it does it perfectly.

This does not solve the problem of political disagreement—people can and will disagree about axioms (assumptions). But it separates the *factual* component of a debate (“does the conclusion follow from the assumptions?”) from the *value* component (“do we accept these assumptions?”).

If machine verification were adopted more broadly in policy analysis, the shape of political discourse would change:

- Debates would center on *assumptions*, not *conclusions*.
- “I don’t trust your analysis” would become “I disagree with Axiom 3.”
- The compiler-as-referee would eliminate an entire category of bad-faith argumentation: the claim that the reasoning itself is flawed.

This is aspirational. But the Cathedral demonstrates that the technology exists. The question is whether the political will exists to use it.

8 Conclusion

The Cathedral’s political significance is not its mathematical content. The Riemann Hypothesis, proved or not, will not change how governments operate. What changes the political landscape is the *demonstration* that:

1. Machine-verified reasoning can be applied to complex systems at significant scale (20,000 lines, 78 files, 644 theorems).
2. Human-AI collaboration can achieve in weeks what institutions require months to accomplish.
3. Every assumption in a complex analysis can be named, tracked, and audited.
4. The tools for this work are publicly available and require no special infrastructure.

The policy question is not whether these capabilities exist—the Cathedral demonstrates that they do. The question is whether democratic institutions will adopt verified reasoning before authoritarian ones do, and whether open science strengthens democratic societies more than it empowers their adversaries.

The Cathedral’s answer—from its security paper, its dual-use paper, and its open methodology—is yes. The asymmetry in verified reasoning favors the defender, the transparency, and the open

society. But this answer must be validated by exactly the people who are responsible for defending those societies.

*The compiler does not choose sides.
But it does distinguish truth from error.
In an age of contested facts,
that may be enough.*

References

- [1] FERC, *Critical Energy Infrastructure Information (CEII) Regulations*, 18 CFR §388.113.
- [2] NERC, *Critical Infrastructure Protection Standards CIP-002 through CIP-014*, 2023.
- [3] NIST, *Artificial Intelligence Risk Management Framework*, AI 100-1, 2023.
- [4] Executive Office of the President, *Executive Order on Safe, Secure, and Trustworthy Artificial Intelligence*, October 2023.
- [5] DARPA, *High-Assurance Cyber Military Systems (HACMS)*, Program Results, 2018.