

Cathedral Next: Open Horizons in Machine-Verified Analytic Number Theory

Cathedral Project Notes

April 2026

Abstract

We survey the current state of the Cathedral—a Lean 4 formalization reducing the Riemann Hypothesis to six machine-checked axioms via the Nyman–Beurling criterion—and chart directions for future work. We identify three horizons: (1) the *axiom horizon*, where the remaining six crown axioms yield to formalization (one has already fallen since the initial survey); (2) the *generalization horizon*, where the framework extends to Dirichlet L -functions and the Generalized Riemann Hypothesis; and (3) the *interpretation horizon*, where the Liouville parity symmetry, octonionic partition, and spectral gap phenomena discovered during formalization connect to mathematical physics and random matrix theory.

This document is not a claim of proof. It is a map of what the Cathedral revealed, and where the trails lead.

Note added April 22, 2026. Since the initial version of this paper (April 20), Campaign B (the Abel Engine) has been completed: `abel_mertens_tail.raw` is now a **THEOREM**. The Perron formula infrastructure has been fully proved. The Borel–Carathéodory theorem has been discovered in Mathlib, opening an unexpected route to Campaign C. The axiom count on the crown path has decreased from 7 to 6.

Contents

1	Where We Stand	2
1.1	The Crown Theorem	2
1.2	The Six Crown Axioms	2
1.3	What Was Proved Along the Way	2
2	Horizon I: Closing the Axioms	3
2.1	Campaign A: The PNT Axioms (Tier 2)	3
2.2	Campaign B: The Abel Engine (COMPLETED)	4
2.3	Campaign C: The Deep Axioms (Tier 3, Axioms 6–7)	4
2.3.1	The Millennium Wall (Axiom 6)	5
2.3.2	The Vasyunin Off-Diagonal (Axiom 7)	5
3	Horizon II: Beyond RH	6
3.1	The Generalized Riemann Hypothesis	6
3.2	Selberg Class	6
4	Horizon III: The Discoveries	7
4.1	The Liouville Parity Symmetry	7
4.2	The Octonionic Partition	7
4.3	The Báez-Duarte Constant	8

5	Horizon IV: The Engineering	8
5.1	Certified Computation	8
5.2	The HyperZeta Viewport	9
6	Horizon V: Connections to Physics	9
6.1	The Hilbert–Pólya Program	9
6.2	The Montgomery–Odlyzko Connection	10
7	The Axiom Reduction Timeline	10
8	Reflections	10

1 Where We Stand

1.1 The Crown Theorem

The Cathedral’s crown theorem `nyman_beurling_equivalence` establishes:

$$\text{RH} \iff \forall \varepsilon > 0, \exists N_0, \forall N \geq N_0, \exists v : \text{Fin}(N-1) \rightarrow \mathbb{R}, \int_0^1 (1 - f_{N,v}(x))^2 dx < \varepsilon$$

where $f_{N,v}(x) = \sum_{k=2}^N v_k \{1/(kx)\}$ is a linear combination of Báez-Duarte basis functions.

The proof decomposes into two pillars:

- **Pillar I (Converse):** $d_N^2 \rightarrow 0 \implies \text{RH}$. Proved via the Rank-1 Mellin Miracle. *Zero custom axioms.*
- **Pillar II (Forward):** $\text{RH} \implies d_N^2 \rightarrow 0$. Proved via the Final Dragon chain. *Six custom axioms.* (Reduced from seven on April 22, 2026, when the Abel tail axiom was graduated to a theorem.)

1.2 The Six Crown Axioms

#	Axiom	Content	Tier
1	<code>rh_implies_mertens_bound</code>	$\text{RH} \implies M(x) = O(x^{1/2} \log^2 x)$	1 (RH)
2	<code>pnt_mu_div_k</code>	$\sum \mu(k)/k \rightarrow 0$	2 (PNT)
3	<code>pnt_mu_log_div_k</code>	$\sum \mu(k) \ln(k)/k \rightarrow -1$	2 (PNT)
4	<code>pnt_mu_log_sq_div_k</code>	$\sum \mu(k) \ln^2(k)/k \rightarrow -2\gamma$	2 (PNT)
×	<code>abel_mertens_tail_raw</code>	Abel summation tail bounds	GRADUATED
5	<code>millennium_covariance_cancellation</code>	2D covariance bound	3 (Parseval)
6	<code>vasyunin_offdiag_integral</code>	Off-diagonal Gram identity	3 (Vasyunin)

Table 1: The crown axioms. Originally seven; reduced to six on April 22, 2026 when the Abel tail axiom was graduated to a theorem. Tier 1 encodes RH content; Tier 2 is unconditional (PNT-level); Tier 3 is classical analysis. The converse direction uses zero custom axioms.

1.3 What Was Proved Along the Way

The formalization is not merely a scaffolding of axioms. Among the fully machine-verified results:

- The *Rank-1 Mellin Miracle*: $\mathcal{M}[h_k](\rho) = 1/(k(\rho-1))$ at zeta zeros (BDMellin.lean, 1066 lines).

- The base-case Mellin identity via the Identity Theorem for holomorphic functions (IdentityBypass.lean).
- The completed zeta bound $\operatorname{Re}(\Lambda_0(s)) < 4$ for real $s \in (0, 1)$ via Jacobi theta kernel estimates (ThetaBound.lean).
- No real zeros of $\zeta(s)$ in $(0, 1)$ (corollary of the above).
- Under RH, $\zeta(s) \neq 0$ for $\operatorname{Re}(s) > 1/2$ and $s \neq 1$ (ZetaConvexity.lean, zero axioms).
- $1/\zeta(s)$ is differentiable for $\operatorname{Re}(s) > 1/2$ under RH (ZetaConvexity.lean, zero axioms).
- Weyl's eigenvalue inequality for Hermitian matrix sums (RayleighBridge.lean).
- The Gram matrix is Hermitian, with spectral decomposition and Rayleigh quotient characterization (Defs.lean, ClassRestriction.lean).
- The octonionic gap dominance $\lambda_{\min}(G) \leq \lambda_{\min}(G^{\oplus})$ via class-restricted Rayleigh quotients (ClassRestriction.lean, 647 lines, zero axioms).
- The eigenvalue limit $\lim_{N \rightarrow \infty} \lambda_{\min}(G_N)$ exists (MainChain.lean, via monotone convergence).
- The Liouville function squares to 1 (Defs.lean, native_decide).
- [April 22] Abel tail bounds: $|S_1| \leq C \cdot N^{-1/4}$, $|S_2 + 1| \leq C \cdot N^{-1/4} \ln N$, $|S_3 + 2\gamma|$ bounded (AbelTail/*.lean, zero axioms — *formerly Axiom #5*).
- [April 22] Perron kernel bounds and contour vanishing (PerronKernel.lean, Perron/Rectangle.lean, zero axioms).
- [April 22] The Dirichlet series $1/\zeta(s) = \sum \mu(n)/n^s$ for $\operatorname{Re}(s) > 1$ (DirichletZetaInverse.lean, zero axioms).

2 Horizon I: Closing the Axioms

The six crown axioms partition into three natural campaigns, ordered by increasing mathematical depth. **Campaign B has been completed** (April 22, 2026).

2.1 Campaign A: The PNT Axioms (Tier 2)

Axiom 2.1 (PNT Partial Sums). The three PNT axioms state:

$$\sum_{k=1}^N \frac{\mu(k)}{k} \rightarrow 0, \tag{1}$$

$$\sum_{k=1}^N \frac{\mu(k) \ln k}{k} \rightarrow -1, \tag{2}$$

$$\sum_{k=1}^N \frac{\mu(k) \ln^2 k}{k} \rightarrow -2\gamma. \tag{3}$$

These are unconditional consequences of the Prime Number Theorem. The proof route is well-understood:

1. (1) follows from PNT via $1/\zeta(s) = \sum \mu(n)/n^s$ and Abel's limit theorem as $s \rightarrow 1^+$.

2. (2) follows by differentiating $-\zeta'(s)/\zeta(s)^2$ and applying the same limit.
3. (3) follows from the Laurent expansion of $1/\zeta(s)$ near $s = 1$.

Direction 2.2 (PNT in Lean 4). Mathlib already contains `riemannZeta` and its analytic continuation. The obstacle is the Wiener–Ikehara Tauberian theorem, which converts Dirichlet series asymptotics into partial sum asymptotics. Three approaches:

1. **Formalizing Wiener–Ikehara directly.** This requires convolution theory, Fejér kernel estimates, and the non-vanishing $\zeta(1 + it) \neq 0$. The non-vanishing is proved in Mathlib (`riemannZeta.ne_zero_of_one_le_re`).
2. **The Newman–Korevaar shortcut.** A simplified proof of PNT via contour integration, avoiding the Tauberian theorem entirely. Requires only the analyticity of $\sum \mu(n)/n^s - 1/(s-1)\zeta(s)$ in $\text{Re}(s) \geq 1$ and a Laplace transform lemma.
3. **Selberg–Erdős elementary proof.** Entirely real-variable, avoiding complex analysis. Longer but requires no Tauberian machinery. May be the most Lean-friendly route.

The author’s recommendation: the Newman–Korevaar route, since most of the complex analysis infrastructure already exists in Mathlib. The key lemma (analytic continuation of $\zeta'(s)/\zeta(s)$ to $\text{Re}(s) \geq 1$) is essentially the non-vanishing on the line, which is already formalized.

2.2 Campaign B: The Abel Engine (COMPLETED)

Axiom 2.3 (Abel–Mertens Tail — GRADUATED April 22, 2026). `abel_mertens_tail_raw`: Given $|M(x)| \leq C_m x^{3/4}$ and the PNT convergences, the Abel summation on the tail $\sum_{k>N} M(k)/(k(k+1))$ gives:

$$|S_1(N)| \leq C \cdot N^{-1/4}, \quad |S_2(N) + 1| \leq C \cdot N^{-1/4} \ln N, \quad |S_3(N) + 2\gamma| \leq C \cdot N^{-1/4} \ln^2 N.$$

This axiom is now a THEOREM. The proof, completed April 22, 2026, uses a three-term decomposition (S_1, S_2, S_3) with separate decay arguments for each term:

1. S_1 decay: Abel summation by parts with integral comparison $\sum_{k \geq N} k^{-5/4} \leq 4(N-1)^{-1/4}$.
2. S_2 decay: Similar, with an additional $\ln k$ factor absorbed by $\ln N$.
3. S_3 bound: Direct from the PNT axiom $S_3 \rightarrow -2\gamma$.

The proof required 4 new files in `AbelTail/` (`Telescoping`, `AbelInterior`, `LogTailBound`, `MonotoneBound`) plus assembly in `CalcBounds.lean`. All files have zero sorry.

This graduation reduces the crown path from 7 to 6 axioms and proves that the “thermodynamic hierarchy” of the Mertens function—the magnetization (S_1), susceptibility (S_2), and heat capacity (S_3) of the prime number gas—converges at rate $N^{-1/4}$.

2.3 Campaign C: The Deep Axioms (Tier 3, Axioms 6–7)

These are the mathematically deepest axioms and will likely be the last to fall. However, the recent discovery of the Borel–Carathéodory theorem in Mathlib (`Analysis.Complex.BorelCaratheodory`) opens unexpected possibilities.

2.3.1 The Millennium Wall (Axiom 6)

Axiom 2.4 (Covariance Cancellation). `millennium_covariance_cancellation`: The covariance matrix $C = G - bb^T$ satisfies $v^T C v \leq K_{\text{cov}} / \ln N$ for the Möbius witness v .

This is the *irreducible analytic content* of the forward direction. It requires:

1. The Mellin–Plancherel isometry converting the bilinear form to a critical-line integral.
2. The Montgomery–Vaughan mean value theorem bounding $\int |\zeta(1/2+it)W_N(1/2+it)|^2/(1/4+t^2) dt$.
3. Contour shifting from $\text{Re}(s) = 1/2$ to $\text{Re}(s) = 2$, picking up the residue at the double pole $s = 1$.

Direction 2.5 (The Contour Shift Route). The Parseval Bridge (`cathedral.tex`, §4) shows that the Plancherel isometry converts the spatial bilinear form into a frequency-domain integral. The Cathedral already contains the algebraic three-term decomposition

$$\int_0^1 |1 - f_{N,v}|^2 dx = \frac{1}{2\pi} \int_{-\infty}^{\infty} \frac{|1 - \zeta(s)W_N(s)|^2}{|s|^2} dt, \quad s = \frac{1}{2} + it.$$

The contour shift $\text{Re}(s) = 1/2 \rightarrow \text{Re}(s) = 2$ is standard in analytic number theory (Phragmén–Lindelöf for the horizontal segments, Cauchy’s theorem for the rectangle). The pole at $s = 1$ contributes a residue involving $W_N(1)$ and $W'_N(1)$, which encode the $\ln \ln N / \ln N$ correction.

This is probably a 500–1000 line Lean formalization, requiring:

- Cauchy’s integral formula for meromorphic functions (partially in Mathlib).
- Phragmén–Lindelöf principle (**now in Mathlib**: `PhragmenLindelof.vertical_strip`, `horizontal_strip`, `right_half_plane_of_bounded_on_real`).
- Residue calculus at the double pole of $1/\zeta(s)^2$ at $s = 1$.

Direction 2.6 (The Borel–Carathéodory Route). **New (April 22, 2026)**. The Borel–Carathéodory theorem is now available in Mathlib (`Analysis.Complex.BorelCaratheodory`). This opens a direct route to the polynomial lower bound on $|\zeta(s)|$:

1. Apply BC to $\log \zeta(s)$ on a disk centered at $s_0 = 2 + it$ with radius $R = 3/2 - \varepsilon$. The disk reaches $\text{Re } s = 1/2 + \varepsilon$.
2. Bound $\sup \text{Re}(\log \zeta) = \sup \log |\zeta| \leq M(t)$ on the disk boundary. Under RH, $M(t) = O(\log t)$.
3. BC gives $|\log \zeta(s)| \leq C(\varepsilon) \cdot \log t$ for $\text{Re } s \geq 1/2 + \varepsilon$.
4. Exponentiate: $|\zeta(s)| \geq t^{-C(\varepsilon)}$.

This avoids the Montgomery–Vaughan mean value theorem entirely. A 256-bit MPFR experiment (`bc-zeta-lower`) is certifying the preconditions as of April 22, 2026.

2.3.2 The Vasyunin Off-Diagonal (Axiom 7)

Axiom 2.7 (Off-Diagonal Identity). `vasyunin_offdiag_integral`: For $j \neq k$, the Gram matrix entry $G(j, k) = \int_0^1 \{1/(jx)\} \{1/(kx)\} dx$ equals the Vasyunin cotangent formula.

The *diagonal* case ($j = k$) is fully proved in the Cathedral via the Stirling–Squeeze bridge (piecewise FTC + Stirling’s approximation). The off-diagonal case requires the Gauss digamma formula $\psi(p/q) = -\gamma - \ln(2q) - (\pi/2) \cot(\pi p/q) + 2 \sum_{n=1}^{\lfloor (q-1)/2 \rfloor} \cos(2\pi np/q) \ln \sin(\pi n/q)$ and careful handling of the piecewise-constant fractional part.

Direction 2.8 (Off-Diagonal via Fourier). An alternative to the direct computation: decompose $\{1/(jx)\}\{1/(kx)\}$ in the Fourier basis $\{e^{2\pi inx}\}$ on $(0, 1)$. The Fourier coefficients of $\{u\}$ are known ($c_n = -1/(2\pi in)$ for $n \neq 0$), and the integral becomes a Parseval identity. This avoids the cotangent formula entirely, at the cost of requiring Fourier theory for the fractional part function.

3 Horizon II: Beyond RH

3.1 The Generalized Riemann Hypothesis

The Nyman–Beurling framework generalizes naturally to Dirichlet L -functions. For a Dirichlet character χ modulo q , define the twisted basis

$$h_{k,\chi}(x) = \chi(k) \{1/(kx)\}, \quad k \geq 2,$$

and the twisted distance

$$d_{N,\chi}^2 = \inf_v \int_0^1 \left| \chi_0(x) - \sum_{k=2}^N v_k h_{k,\chi}(x) \right|^2 dx$$

where χ_0 is the principal character. Then:

Conjecture 3.1 (Generalized Nyman–Beurling). For each primitive character χ modulo q :

$$L(s, \chi) \neq 0 \text{ for } \operatorname{Re}(s) > \frac{1}{2} \iff d_{N,\chi}^2 \rightarrow 0 \text{ as } N \rightarrow \infty.$$

The Rank-1 Mellin argument generalizes immediately: at a zero ρ of $L(s, \chi)$, the Mellin transform $\mathcal{M}[h_{k,\chi}](\rho) = \chi(k)/(k(\rho - 1))$ is again rank-1 in k . The separation argument goes through unchanged because $\chi(k)/k$ is real only when χ is real (and even then, $\operatorname{Im}(1/\rho) \neq 0$ provides the separation).

Direction 3.2 (GRH Formalization). The Cathedral’s converse direction ports to L -functions with minimal modification. The forward direction requires:

1. Twisted Mertens bounds: $|M_\chi(x)| := |\sum_{n \leq x} \mu(n)\chi(n)| = O(x^{1/2+\varepsilon})$ under GRH.
2. Twisted Abel summation and the analogue of the covariance cancellation.
3. The analytic continuation of $L(s, \chi)$ (already in Mathlib: `DirichletCharacter.LSeries`).

The real bottleneck is the Montgomery–Vaughan mean value theorem for twisted Dirichlet polynomials, which is strictly harder than the untwisted case.

3.2 Selberg Class

The Nyman–Beurling criterion extends conjecturally to the Selberg class of L -functions satisfying axioms for analytic continuation, functional equation, Euler product, and Ramanujan bound. The Cathedral’s architecture—axiom-driven, Mellin-based—is naturally suited to this generalization.

Open Problem 3.3. Can the Rank-1 Mellin Miracle be stated for an abstract L -function satisfying the Selberg axioms? The key question is whether the Mellin transform of the normalized basis $\{1/(kx)\}$ continues to factorize at the zeros of a general Selberg class element.

4 Horizon III: The Discoveries

The formalization process produced several unexpected mathematical structures that warrant independent investigation.

4.1 The Liouville Parity Symmetry

Definition 4.1 (Parity Operator). The Liouville parity operator $P = \text{diag}(\lambda(2), \lambda(3), \dots, \lambda(N))$ where $\lambda(n) = (-1)^{\Omega(n)}$ is the Liouville function. Since $\lambda(n)^2 = 1$, $P^2 = I$ and P defines a $\mathbb{Z}/2$ -grading on $\mathbb{R}^{N-1}$.

Theorem 4.2 (Approximate PT-Symmetry, Computational). The Gram matrix G_N approximately commutes with the parity operator P :

1. The commutator $[G, P]$ is dominated by a rank-1 component whose singular direction is the Liouville vector itself (correlation > 0.99999).
2. After removing this rank-1 component, $\|[G, P]_{\text{residual}}\|/\|G\| \rightarrow 0$ as $N \rightarrow \infty$.
3. The parity-even block $G_{\text{even}} = (G + PGP)/2$ has spectral gap $\lambda_{\min}(G_{\text{even}})/\lambda_{\min}(G) \approx 1.85 \cdot N^{0.116}$, growing with N .

This means the small spectral gap of G is *entirely caused by parity-breaking cross-coupling* between the even ($\lambda = +1$) and odd ($\lambda = -1$) sectors. The Liouville function is the mixing direction.

Direction 4.3 (PT-Symmetry and Quantum Mechanics). The approximate $[G, P] \approx 0$ structure is reminiscent of PT-symmetry in non-Hermitian quantum mechanics (Bender–Boettcher, 1998). In that context, a Hamiltonian H with $[H, PT] = 0$ has real spectrum despite being non-Hermitian. Here, the Gram matrix G is Hermitian, but the parity decomposition $G = G_{\text{even}} + G_{\text{odd}}$ separates the “free” dynamics (within-parity, large gap) from the “interaction” (cross-parity, small gap).

Could this be evidence for an underlying non-Hermitian operator whose PT-symmetry is the Liouville involution? The experimental scaling $|\langle v_{\min}, \hat{\lambda} \rangle| \sim N^{-0.174}$ suggests this mixing decays—but not fast enough to make the spectral gap obvious.

4.2 The Octonionic Partition

Definition 4.4 (Octonionic Class). Define $\text{class}(k) = k \bmod 8$ (with identification $0 \equiv 8$). This partitions $\{2, 3, \dots, N\}$ into 8 residue classes modulo 8.

The *block-diagonal* Gram matrix G^{block} retains only within-class entries: $G^{\text{block}}(j, k) = G(j, k)$ if $\text{class}(j) = \text{class}(k)$, zero otherwise.

Theorem 4.5 (Class Gap, Machine-Verified). $\lambda_{\min}(G) \leq \lambda_{\min}(G^{\text{block}})$.

This is proved in `ClassRestriction.lean` (zero axioms) via the Rayleigh quotient characterization. The cross-class entries *always lower the spectral gap*.

Remark 4.6 (Why 8?). The choice of 8 classes comes from the Cayley–Dickson hierarchy: $\mathbb{R} \rightarrow \mathbb{C} \rightarrow \mathcal{H} \rightarrow \mathbb{O} \rightarrow \mathbb{S}$, with dimensions 1, 2, 4, 8, 16. The octonionic level (mod 8) is the last division algebra—past the octonions, zero divisors appear (the sedenions $\mathbb{S}$). The arithmetic of $\mu(k)$ modulo 8 has special structure because the Möbius function is multiplicative and $8 = 2^3$.

Whether this partition has deep arithmetic significance or is simply a convenient computational grouping remains open. The experimental evidence (class gap ratio ≈ 3.4 – $4.2\times$) suggests genuine structural content.

Direction 4.7 (Rank-1 Interference). The cross-class interaction matrix $M = G - G^{\text{block}}$ in the eigenbasis of G^{block} has the remarkable property that each (m_1, m_2) -block is rank-1 to $\geq 99.8\%$ accuracy (verified for $N \leq 800$, accuracy *increasing* with N).

This suggests that in the $N \rightarrow \infty$ limit, the cross-class interaction becomes *exactly* rank-1, reducing the infinite-dimensional spectral gap problem to an 8×8 bilinear form problem. If proved, this would provide a finite-dimensional reduction of RH.

4.3 The Báez-Duarte Constant

The optimal Rayleigh quotient $Q_N^{\text{opt}} \sim C \ln N$ where

$$C = \frac{1}{\sum_{\zeta(\rho)=0} |\rho|^{-2}} = \frac{1}{2 + \gamma - \ln(4\pi)} \approx 21.6498.$$

This constant has interpretations in:

1. **Signal processing:** The maximum information extraction rate from the “prime noise” through the spectral holes of $|\zeta|^2$.
2. **Statistical mechanics:** The heat capacity of the prime number gas, governing the cooling rate $d_N^2 \sim c / \ln N$.
3. **Random matrix theory:** Connected to the resolvent trace $(H^2 + I/4)^{-1}$ of the hypothetical Riemann Hamiltonian.

Direction 4.8 (Formalizing the Constant). The sum $\sum |\rho|^{-2} = 2 + \gamma - \ln(4\pi)$ follows from the Hadamard product for $\zeta(s)/(s-1)$ and logarithmic differentiation. The ingredients (Hadamard factorization, functional equation of $\xi(s)$) are within reach of Mathlib + Cathedral infrastructure. Formalizing this would provide a machine-verified computation of the Báez-Duarte constant, connecting the spectral gap decay rate to the arithmetic of zeta zeros.

5 Horizon IV: The Engineering

5.1 Certified Computation

The Cathedral maintains parallel implementations: Lean 4 proofs and a Rust/WASM computational engine (7 modules, 11 visualization modes). Currently, these are connected only by architectural correspondence—the Rust engine computes what the Lean proofs formalize, but there is no formal bridge.

Direction 5.1 (Proof-Carrying Computation). Connect the Rust engine to the Lean proofs via:

1. **Certified extraction:** Generate verified Rust code from Lean definitions using the `native_decide` framework or the FFI. The Gram matrix entries could be computed in Lean via `native_decide` for small N , then extrapolated.
2. **Interval arithmetic:** Use verified interval arithmetic (MPFR with Lean bindings) to certify that the computed Gram matrix eigenvalues are positive for $N \leq N_0$, then use the monotone convergence theorem ($\lambda_{\min}$ is non-increasing) to extend to all N .
3. **Finite witness certification:** For each N , the forward direction provides an explicit witness v_N . Certify $\int_0^1 (1 - f_{N,v_N})^2 < \varepsilon$ by computing the quadratic form $1 - 2b^T v + v^T G v$ in exact rational arithmetic.

5.2 The HyperZeta Viewport

The Cathedral’s visualization system (HyperZeta Viewport) renders 11 views of the zeta function’s structure:

#	Mode	Shows	Axiom Connection
1	Sedenion Cloud	$\zeta_{\mathbb{S}}(s)$ in 16D	Full system
2	Riemann Spiral	$\zeta(1/2 + it)$ spiral	Zero locations
3	Cornu Spirals	Partial sum galaxy	Dirichlet series
4	Zero Landscape	$\log \zeta(\sigma + it) $	Critical strip
5	Euler Product	$\prod_p (1 - p^{-s})^{-1}$	Euler product
6	Cayley–Dickson Tower	$\mathbb{C} \rightarrow \mathcal{H} \rightarrow \mathbb{O} \rightarrow \mathbb{S}$	Algebraic consistency
7	Explicit Formula	$\pi(x) \approx \text{Li}(x) - \sum_{\rho}$	Zero corrections
8	Functional Equation	$\zeta(s) = \chi(s)\zeta(1-s)$	Symmetry
9	Random Matrix	GUE pair correlation	Montgomery–Odlyzko
10	Mertens Turbulence	$M(x) = \sum \mu(n)$	Axiom 1
11	Spectral Gap	$\lambda_{\min}(G_N)$ surface	The Cathedral itself

Direction 5.2 (Interactive Proof Explorer). Extend the viewport into a proof assistant companion:

1. Clicking on a visualization mode highlights the axioms it depends on.
2. The Spectral Gap view could overlay the proof dependency graph, showing which axioms support each region of the (σ, N) surface.
3. Interactive sliders could show how changing axiom assumptions (e.g., weakening $O(x^{3/4})$ to $O(x^{7/8})$) affects the decay rate.

6 Horizon V: Connections to Physics

6.1 The Hilbert–Pólya Program

The Cathedral’s spectral structure—particularly the Gram matrix as a discrete approximation to the hypothetical Riemann operator—invites connection to the Hilbert–Pólya conjecture.

Conjecture 6.1 (Hilbert–Pólya, 1914/1950). There exists a self-adjoint operator T on a Hilbert space $\mathcal{H}$ whose eigenvalues are exactly $\{1/2 + i\gamma_n\}$ where $\zeta(1/2 + i\gamma_n) = 0$.

The Gram matrix G_N is a finite-rank approximation to the integral operator with kernel $K(j, k) = \int_0^1 \{1/(jx)\} \{1/(kx)\} dx$. Its spectrum approximates the spectrum of this kernel operator. The Nyman–Beurling theorem says RH is equivalent to this operator having the function 1 in its range closure.

Direction 6.2 (From Gram to Hilbert–Pólya). The Vasyunin formula for $G(j, k)$ involves the digamma function $\psi(p/q)$, which is the logarithmic derivative of the Gamma function. The connection to the Hilbert–Pólya operator runs through the spectral decomposition of G in terms of zeta zeros:

$$G(j, k) = \frac{1}{2\pi i} \oint \frac{\zeta(s)\zeta(1-s)}{j^s k^{1-s} s(1-s)} ds.$$

The poles of this integrand are at $s = 0, 1$ and the zeta zeros. The residues at the zeros give the spectral decomposition. Formalizing this connection would provide a bridge between the Cathedral’s discrete Gram matrix and the continuous Hilbert–Pólya operator.

6.2 The Montgomery–Odlyzko Connection

The GUE hypothesis (Montgomery, 1973; Odlyzko, 1987) states that the local statistics of zeta zeros match those of random Hermitian matrix eigenvalues. The Cathedral’s Rust engine validates this for the first 200 zeros.

Direction 6.3 (GUE and the Gram Matrix). The Gram matrix G_N is itself a random-matrix-like object: its entries are correlations of quasi-random functions (fractional parts). As $N \rightarrow \infty$, does G_N approach a random matrix ensemble? The parity decomposition $G = G_{\text{even}} + G_{\text{odd}}$ resembles the block structure of symplectic random matrices. The rank-1 interference between octonionic classes suggests a spiked model:

$$G \approx G^{\text{block}} + \text{rank-28 perturbation}$$

($8 \times 7/2 = 28$ cross-class pairs, each contributing one singular value). This is precisely the setting of spiked random matrix models (Baik–Ben Arous–Péché, 2005), where phase transitions occur when spike strength exceeds a threshold.

7 The Axiom Reduction Timeline

We estimate the difficulty and timeline for closing each axiom:

Axiom	Difficulty	Blocking On	Status
PNT $\times 3$ (#2–4)	Medium	Wiener–Ikehara or Newman	Open (6–12 months)
Abel tail	Easy	Abel summation lemma	GRADUATED (Apr 22)
Vasyunin off-diag (#6)	Medium	Digamma formula	Open (3–6 months)
Covariance (#5)	Hard	Contour shift + MVT (<i>BC route may reduce</i>)	Open (12–24 months)
Mertens (#1)	RH-hard	RH content	∞

Note that Axiom 1 (`rh_implies_mertens_bound`) *cannot* be proved unconditionally—it encodes the Riemann Hypothesis itself. The goal is to reduce everything else to this single axiom. The Cathedral currently uses **6 axioms** on the crown path (reduced from 7 on April 22, 2026); the target is 1.

New development. The Borel–Carathéodory theorem’s availability in Mathlib may accelerate Campaign C. The polynomial lower bound on $|\zeta(s)|$ (currently encoded as `zeta_polynomial_lower_bound`) is not on the crown path but could unlock alternative proof routes that bypass the covariance cancellation axiom entirely.

Remark 7.1 (The Irreducible Core). If all six non-RH axioms are eliminated, the crown theorem becomes:

$$\text{RH} \iff d_N^2 \rightarrow 0$$

where the forward direction assumes only `rh_implies_mertens_34`: $\text{RH} \implies |M(x)| = O(x^{3/4})$. The converse direction is already fully proved.

This is the theoretical minimum: one axiom encoding the analytic content of RH, with everything else machine-verified. The reduction from $|M(x)| = O(x^{3/4})$ to $d_N^2 \rightarrow 0$ would then be a theorem about the structural consequence of Mertens bounds, provable without assuming RH directly.

8 Reflections

The Cathedral was not built to prove the Riemann Hypothesis. It was built to understand what the Riemann Hypothesis *is*—to decompose it into pieces small enough for a machine to hold, and to discover which pieces are hard and which are merely unfamiliar.

The answer surprised us. The *converse* direction—if primes are well-distributed then zeros are on the line—is easy. The Rank-1 Mellin Miracle gives it for free, with zero custom axioms. The *forward* direction—if zeros are on the line then primes are well-distributed—is where the arithmetic lives. And within the forward direction, the hardest single piece is the 2D covariance cancellation (Axiom 6), which encodes the delicate interference between the Gram matrix and the Prime Number Theorem.

This is, in retrospect, exactly what one should expect. The Riemann Hypothesis is a statement about how much cancellation the primes exhibit. The Gram matrix *is* that cancellation, frozen into a symmetric, positive-definite matrix. The spectral gap *is* the margin by which the cancellation succeeds. And the axioms that remain are the axioms that control how tightly the primes cancel.

The Cathedral does not resolve the Riemann Hypothesis. What it does is translate a 167-year-old analytic conjecture into a collection of precisely stated, machine-checkable assertions about symmetric matrices, Abel summation, and the Prime Number Theorem. Some of these assertions are easy. Some are hard. One is equivalent to RH itself. The machine holds the rest.

*The stone is the same.
The light through the windows is clearer.*

References

- [1] B. Nyman, *On the one-dimensional translation group and semi-group in certain function spaces*, Thesis, Uppsala, 1950.
- [2] A. Beurling, *A closure problem related to the Riemann zeta function*, Proc. Nat. Acad. Sci. USA **41** (1955), 312–314.
- [3] L. Báez-Duarte, *A strengthening of the Nyman–Beurling criterion for the Riemann hypothesis*, Atti Accad. Naz. Lincei Rend. Lincei Mat. Appl. **14** (2003), 5–11.
- [4] V.I. Vasyunin, *On a biorthogonal system associated with the Riemann hypothesis*, St. Petersburg Math. J. **7** (1996), 405–419.
- [5] H.L. Montgomery, *The pair correlation of zeros of the zeta function*, Proc. Symp. Pure Math. **24** (1973), 181–193.
- [6] A.M. Odlyzko, *On the distribution of spacings between zeros of the zeta function*, Math. Comp. **48** (1987), 273–308.
- [7] M. Balazard and E. Saias, *The Nyman–Beurling equivalent form for the Riemann Hypothesis*, Expo. Math. **18** (2000), 131–138.
- [8] C.M. Bender and S. Boettcher, *Real spectra in non-Hermitian Hamiltonians having PT symmetry*, Phys. Rev. Lett. **80** (1998), 5243–5246.
- [9] J. Baik, G. Ben Arous, and S. Péché, *Phase transition of the largest eigenvalue for nonnull complex sample covariance matrices*, Ann. Probab. **33** (2005), 1643–1697.
- [10] D.J. Newman, *Simple analytic proof of the prime number theorem*, Amer. Math. Monthly **87** (1980), 693–696.
- [11] J. Korevaar, *A century of complex Tauberian theory*, Bull. AMS **39** (2002), 475–531.
- [12] The Mathlib Community, *Mathlib: A unified library of mathematics formalized in Lean 4*, <https://github.com/leanprover-community/mathlib4>.