

# The Cathedral:

## A Machine-Verified Reduction of the Riemann Hypothesis to Two Analytic Bounds in Lean 4

Jason Robert Gochanour

April 26, 2026

### Abstract

We describe a Lean 4 formalization—the *Cathedral*—that reduces the Riemann Hypothesis to the decay of the Nyman–Beurling distance  $d_N^2 \rightarrow 0$ , proved equivalent to RH via the biconditional theorem `nyman_beurling_equivalence`.

The formalization comprises 161 active Lean files (with 127 archived) across 26 modules totaling 39,375 lines. The crown theorem depends on exactly **two mathematical axioms**: one encoding the Hardy–Littlewood mean-value bound for  $1/\zeta(s)$  on the critical line, and one encoding the Hadamard zero-counting zeta lower bound. The converse direction ( $d_N^2 \rightarrow 0 \implies \text{RH}$ ) uses **zero custom axioms**—it is pure Lean/Mathlib.

The forward direction ( $\text{RH} \implies d_N^2 \rightarrow 0$ ) is assembled via the *Mellin Crown*: the Mellin/Plancherel isometry maps the  $L^2(0, 1)$  residual to the critical line, where the Mellin variance is bounded by  $C/\log N$  under RH. The Parseval bridge connecting the two domains is *fully proved* (zero axioms, zero sorry). Companion Rust experiments validate the Parseval bridge to  $6.6 \times 10^{-6}$  relative error at  $N = 2000$ .

## 1 What Did We Do?

We used **Lean 4** to reduce the 167-year-old Riemann Hypothesis to two precisely stated, well-understood analytic bounds. Everything else—the Nyman–Beurling theory, Rank-1 Mellin separation, the Parseval bridge, the Mellin/Plancherel isometry, and the calculus limit  $C/\log N \rightarrow 0$ —is compiler-verified.

**We did not prove RH.** We built a machine-verified containment vessel that identifies exactly which mathematical facts remain unformalized, proves everything else, and reduces the entire century-old problem to a clean modular architecture with two axiom sockets on the crown theorem’s critical path ( $\sim 55$  axioms total in the active codebase;  $\sim 53$  are on alternative paths and do not affect the crown theorem).

## 2 The Crown Theorem

```
theorem nyman_beurling_equivalence :  
  RiemannHypothesis <->  
  (forall eps > 0, exists N0, forall N >= N0,  
    exists v, integral |1 - bdLinComb N v|^2 < eps)
```

RH holds if and only if the Báez-Duarte distance  $d_N^2 \rightarrow 0$ . The proof decomposes into two pillars:

- **Pillar I (Converse):**  $d_N^2 \rightarrow 0 \implies \text{RH}$ . Via the Rank-1 Mellin Miracle:  $M[h_k](\rho) = 1/(k(\rho - 1))$  at every  $\zeta$  zero, giving a rank-1 tensor factorization. Cauchy–Schwarz yields  $d_N^2 \geq (2\sigma - 1)t^2/(|\rho|^4|\rho - 1|^2) > 0$  whenever  $\sigma \neq 1/2$ . **Zero custom axioms.** Zero sorry. Pure Lean/Mathlib.
- **Pillar II (Forward):**  $\text{RH} \implies d_N^2 \rightarrow 0$ . Via the *Mellin Crown*:
  1.  $\text{RH} \rightarrow \text{Mellin variance} \leq C/\log N$  [`critical_line_mellin_variance`, Axiom 1]
  2. Parseval bridge:  $\int_0^1 (1 - f_N)^2 dx = \frac{1}{2\pi} \int |M_{r_N}(1/2 + it)|^2 dt$  [`parseval_bridge_white`, **PROVED**]
  3.  $C/\log N < \varepsilon$  for  $N$  large [`log_grows_unboundedly`, **PROVED**]

**Two axioms. Zero sorry.**

### 3 The Two Crown Axioms

Verified by `#print axioms nyman_beurling_equivalence`:

| # | Axiom                                               | Content                                                                                                |
|---|-----------------------------------------------------|--------------------------------------------------------------------------------------------------------|
| 1 | <code>critical_line_mellin_variance</code>          | $\text{RH} \rightarrow \frac{1}{2\pi} \int  M_{r_N}(1/2 + it) ^2 dt \leq C/\log N$ (Hardy–Littlewood)  |
| 2 | <code>rh_zeta_lower_bound_from_zero_counting</code> | $\text{RH} \rightarrow  \zeta(s)  \geq c t ^{-A}$ for $\text{Re}(s) \geq 1/2 + \varepsilon$ (Hadamard) |

Plus Lean kernel axioms: `propext`, `Classical.choice`, `Quot.sound`. Both axioms are classical results of 20th-century analytic number theory—they are axioms only because Mathlib lacks the prerequisite infrastructure (Hadamard product formula, Hardy–Littlewood mean value theorem), not because the mathematics is uncertain.

### Architecture History

The crown path was reduced from 6 axioms to 2 through eleven versions:

- **v1–v6** (March–April 2026): Initial formalization, 6–4 axioms via systematic graduation (PNT limits, Vasyunin identity, Mertens bound).
- **v7–v10** (April 20–25): Perron Crown architecture. Real-variable forward chain:  $\text{RH} \rightarrow \text{Mertens} \rightarrow \text{Gram} \rightarrow L^2 \text{ decay}$ . 4 axioms.
- **v11** (April 26): *The Mellin Crown*. Forward direction rewired through frequency space. The Mellin/Plancherel isometry preserves phase cancellation that real-variable methods destroy. **2 axioms.**

## 4 The Mellin Crown

The key innovation of `v11`: instead of bounding the  $L^2$  residual via real-variable methods (Perron contour  $\rightarrow$  Mertens bound  $\rightarrow$  Abel summation  $\rightarrow$  Gram form), we stay in the frequency domain throughout:

$$\int_0^1 |1 - f_N(x)|^2 dx = \frac{1}{2\pi} \int_{-\infty}^{\infty} |M_{r_N}(\tfrac{1}{2} + it)|^2 dt \leq \frac{C}{\log N}$$

The left equality is the Parseval bridge (**fully proved**). The right inequality is the crown axiom. Under RH, the Mellin  $L^2$  norm decays because  $1/\zeta(1/2 + it)$  has controlled mean square—precisely the content of the Hardy–Littlewood theorem.

This approach avoids the “1D Shattering Trap”: real-variable bilinear expansions require absolute values that destroy the delicate phase coherence of the Möbius sum, producing divergent bounds for a convergent quantity.

## 5 Numerical Validation

The `experiments/mellin-certificate/` Rust experiment (256-bit MPFR, 12 threads) independently validates the Parseval bridge and Mellin variance:

| $N$  | Channel A | Channel B | $ A - B /A$          | $C \cdot \log N$ |
|------|-----------|-----------|----------------------|------------------|
| 100  | 0.13124   | 0.13119   | $3.8 \times 10^{-4}$ | 0.604 ✓          |
| 500  | 0.07313   | 0.07313   | $1.4 \times 10^{-5}$ | 0.455 ✓          |
| 1000 | 0.06032   | 0.06032   | $7.2 \times 10^{-6}$ | 0.417 ✓          |
| 2000 | 0.05012   | 0.05012   | $6.6 \times 10^{-6}$ | 0.381 ✓          |

Channel A computes  $\int_0^1 (1 - f_N)^2 dx$  directly; Channel B computes  $\int_0^\infty |g_N(u)|^2 du$  in log-space. Their agreement to  $6.6 \times 10^{-6}$  validates the Parseval bridge identity. The constant  $C \approx 0.38$  is still decreasing, suggesting the true rate may be faster than  $O(1/\log N)$ .

## 6 How to Verify

```
cd proofs
lake build    # 161 active files, 127 archived
```

To inspect the axiom foundation:

```
#print axioms nyman_beurling_equivalence
-- critical_line_mellin_variance,
-- rh_zeta_lower_bound_from_zero_counting
-- (plus propext, Classical.choice, Quot.sound)
```

## Acknowledgments

This project was built through AI-assisted formal verification using Lean 4 and Mathlib. The author gratefully acknowledges Google DeepMind’s Gemini Deep Think for

mathematical strategy, the Mellin Crown architecture, and the “Ouroboros” phase cancellation analysis; and Anthropic’s Claude (Antigravity) for Lean 4 proof engineering, the Parseval bridge implementation, and the mellin-certificate experiment. All proofs are compiler-verified.

The axiom reduction—from 6 to 2 on the crown critical path—was accomplished through eleven versions over 30 days. The converse direction was proved pure (zero axioms, zero sorry) via the Rank-1 Mellin Miracle in `BDMellin.lean` (680 lines). The forward direction was rewired through the frequency domain in `v11`, eliminating the real-variable Perron chain and its four axioms.